

CS231n Convolutional Neural Networks for Visual Recognition

In the previous assignment, you implemented and trained your own ConvNets. In this assignment, we will explore many of the ideas we have discussed in lectures. Specifically, you will:

- Reduce overfitting using **dropout** and **data augmentation**
- Combine models into **ensembles** to boost performance
- Use **transfer learning** to adapt a pretrained model to a new dataset
- Use data gradients to visualize **saliency maps** and create **fooling images**

Setup

You can work on the assignment in one of two ways: locally on your own machine, or on a virtual machine through [Terminal](#).

Working locally

Get the code: Download the starter code [here](#).

[Optional] virtual environment: Once you have unzipped the starter code, you might want to create a [virtual environment](#) for the project. If you choose not to use a virtual environment, it is up to you to make sure that all dependencies for the code are installed on your machine. To set up a virtual environment, run the following:

```
cd assignment3
sudo pip install virtualenv      # This may already be installed
virtualenv .env                  # Create a virtual environment
source .env/bin/activate         # Activate the virtual environment
pip install -r requirements.txt  # Install dependencies
# Work on the assignment for a while ...
deactivate                       # Exit the virtual environment
```

You can reuse the virtual environment that you created for the first or second assignment, but you will need to run `pip install -r requirements.txt` after activating it to install additional dependencies required by this assignment.

Download data: Once you have the starter code, you will need to download the CIFAR-10 dataset, the TinyImageNet-100-A and TinyImageNet-100-B datasets, and pretrained models for the TinyImageNet-100-A dataset.

Run the following from the `assignment3` directory:

NOTE: After downloading and unpacking, the data and pretrained models will take about 900MB of disk space.

```
cd cs231n/datasets
./get_datasets.sh
./get_tiny_imagenet_splits.sh
./get_pretrained_models.sh
```

Compile the Cython extension: Convolutional Neural Networks require a very efficient implementation. We have implemented of the functionality using [Cython](#); you will need to compile the Cython extension before you can run the code. From the `cs231n` directory, run the following command:

```
python setup.py build_ext --inplace
```

Start IPython: After you have downloaded the data and compiled the Cython extensions, you should start the IPython notebook server from the `assignment3` directory. If you are unfamiliar with IPython, you should read our [IPython tutorial](#).

Working on Terminal

We have created a Terminal snapshot that is preconfigured for this assignment; you can [find it here](#). Terminal allows you to work on the assignment from your browser. You can find a tutorial on how to use it [here](#).

Submitting your work:

Whether you work on the assignment locally or using Terminal, once you are done working run the `collectSubmission.sh` script; this will produce a file called `assignment3.zip`. Upload this file to your dropbox on [the coursework](#) page for the course.

Q1: Dropout and Data Augmentation

In the IPython notebook `q1.ipynb` you will implement dropout and several forms of data augmentation. This will allow you to reduce overfitting when training on a small subset of the CIFAR-10 dataset.

Q2: TinyImageNet and Model Ensembles

In the IPython notebook `q2.ipynb` we will introduce the TinyImageNet-100-A dataset. You will try to classify examples from this dataset by hand, and you will combine pretrained models into an ensemble to boost your performance on this dataset.

Q3: Transfer Learning

In the IPython notebook `q3.ipynb` you will implement several forms of transfer learning. You will use adapt a pretrained model for TinyImageNet-100-A to achieve reasonable performance with minimal training time on the TinyImageNet-100-B dataset.

Q4: Visualizing and Breaking ConvNets

In the IPython notebook `q4.ipynb` you will visualize data gradients and you will generate images to fool a pretrained ConvNet.

 [cs231n](#)

 [cs231n](#)

karpathy@cs.stanford.edu